

See discussions, stats, and author profiles for this publication at: <https://www.researchgate.net/publication/329512940>

CUDA Tutorial – How to Start with CUDA?

Presentation · December 2018

DOI: 10.13140/RG.2.2.22460.49289

CITATIONS

0

READS

2,132

1 author:



Ali Tourani

University of Luxembourg

25 PUBLICATIONS 78 CITATIONS

[SEE PROFILE](#)

Some of the authors of this publication are also working on these related projects:



Search Engine [View project](#)



Deep-Learning Applications in Intelligent Transportation Systems [View project](#)



CUDA Tutorial



ALI TOURANI

MSC. SOFTWARE ENGINEERING

A.TOURANI1991@GMAIL.COM

Scopes

- ▶ Introduction
- ▶ CUDA key concepts
- ▶ CUDA threads
- ▶ CUDA performance
- ▶ CUDA memories
- ▶ CUDA installation
- ▶ CUDA matrix multiplication

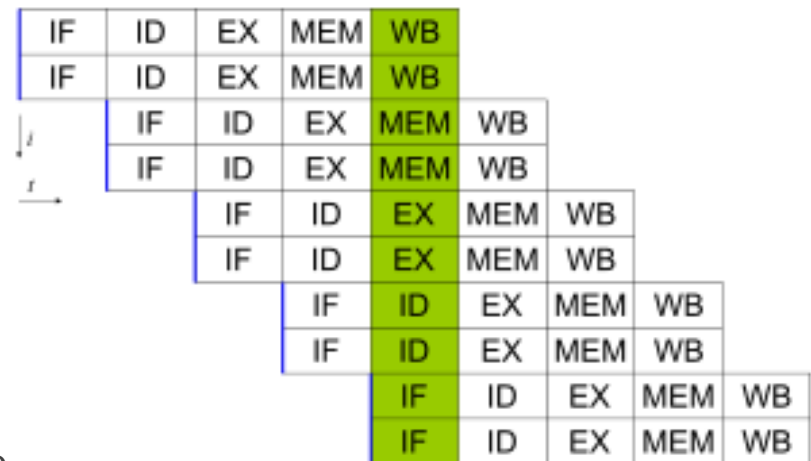


Introduction

- ▶ CUDA: a parallel computing platform and API model
 - ▶ Developed by NVidia
- ▶ Utilization of the power of NVidia GPUs
 - ▶ Doing graphical calculations
 - ▶ Perform general computing tasks
 - ▶ Image Processing
 - ▶ Deep Learning
 - ▶ Multiplying matrices
- ▶ Requirements:
 - ▶ C/C++/Python programming and a NVidia GPU card!

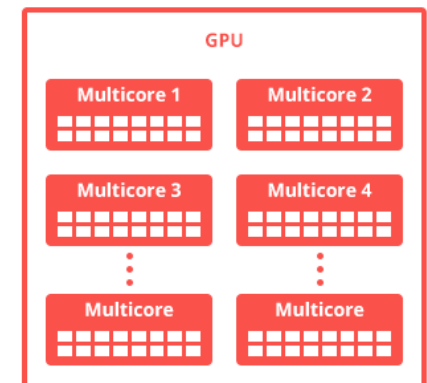
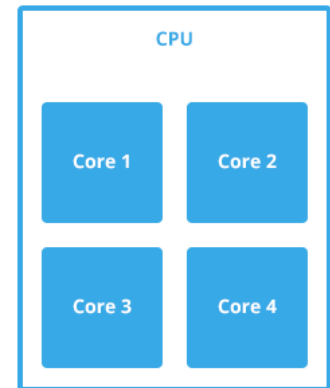
Introduction

- ▶ Parallelism in the CPU
 - ▶ Instruction fetch (IF)
 - ▶ Instruction decode (ID)
 - ▶ Instruction execute (EX)
 - ▶ Memory access (MEM)
 - ▶ Register write-back (WB)
- ▶ Pipelining
 - ▶ Instruction Level Parallelism (ILP)

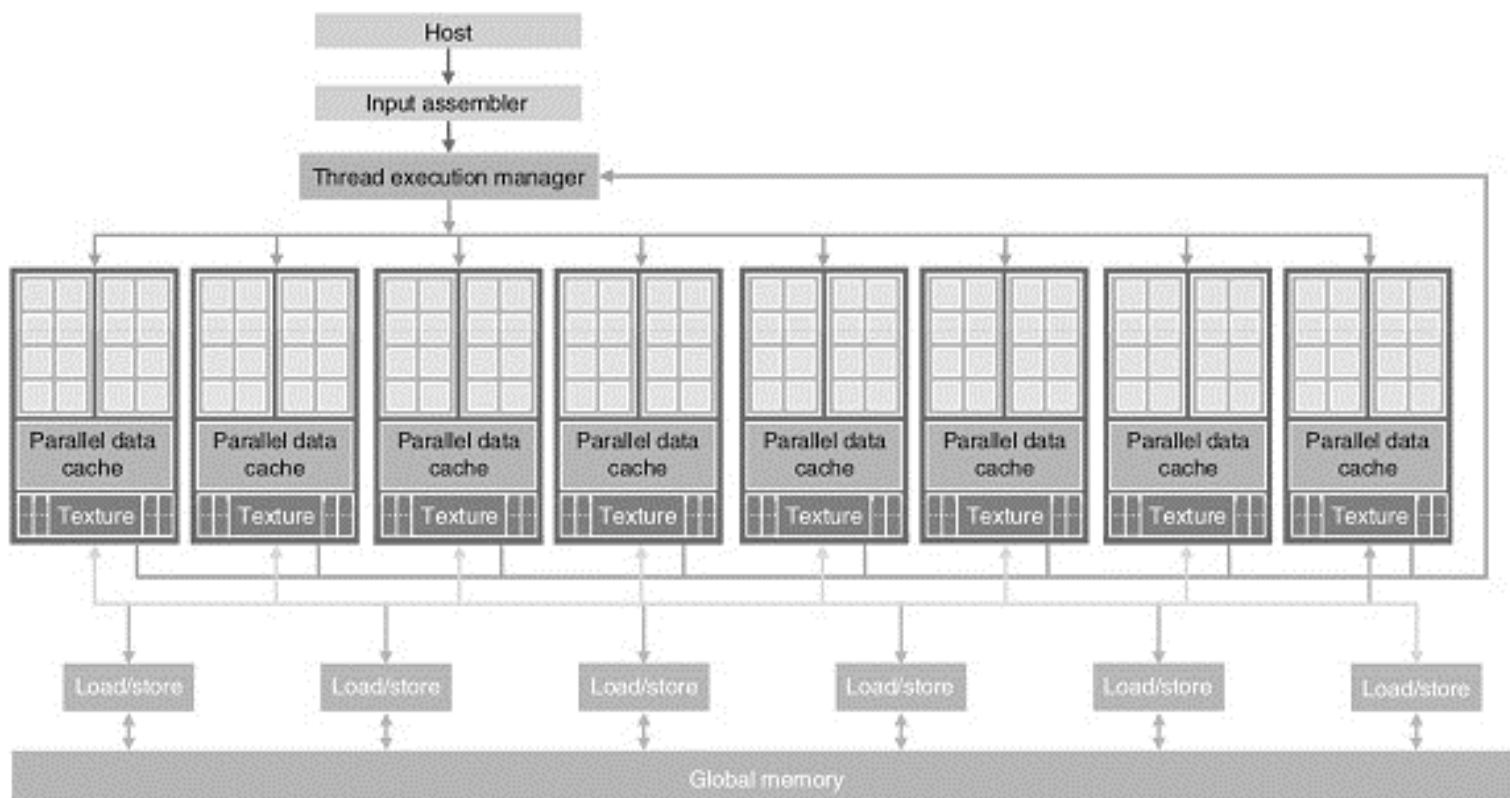


Introduction

- ▶ Parallelism in the GPU
 - ▶ Many-core processors
 - ▶ Operation of large chunks of data
 - ▶ Massively-parallel programs
 - ▶ Efficient for SPMD (Single Program, Multiple Data)
 - ▶ No virtual memory
 - ▶ No interrupts

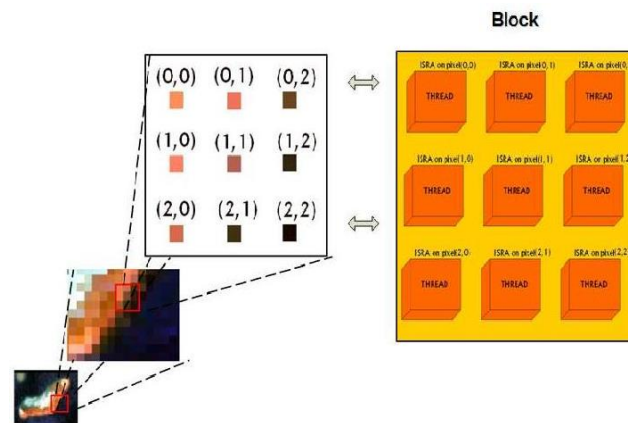
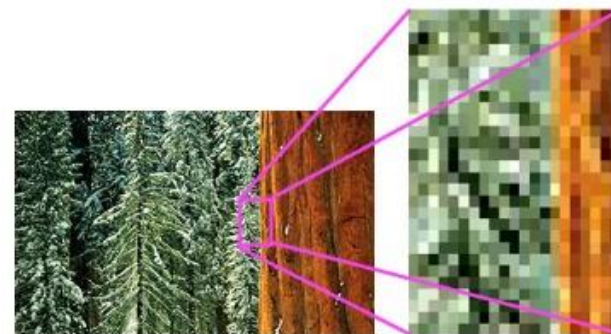


Introduction



CUDA key concepts

- ▶ Data Parallelism in CUDA
 - ▶ Computationally intensive applications
 - ▶ Such as rendering pixels
 - ▶ Threads (workers) as the main tools
 - ▶ Pixels to threads mapping $\rightarrow O(1)$
 - ▶ Each thread processes one pixel



CUDA key concepts

► Structure of CUDA (for a C program)

► NVCC (NVIDIA C Compiler)

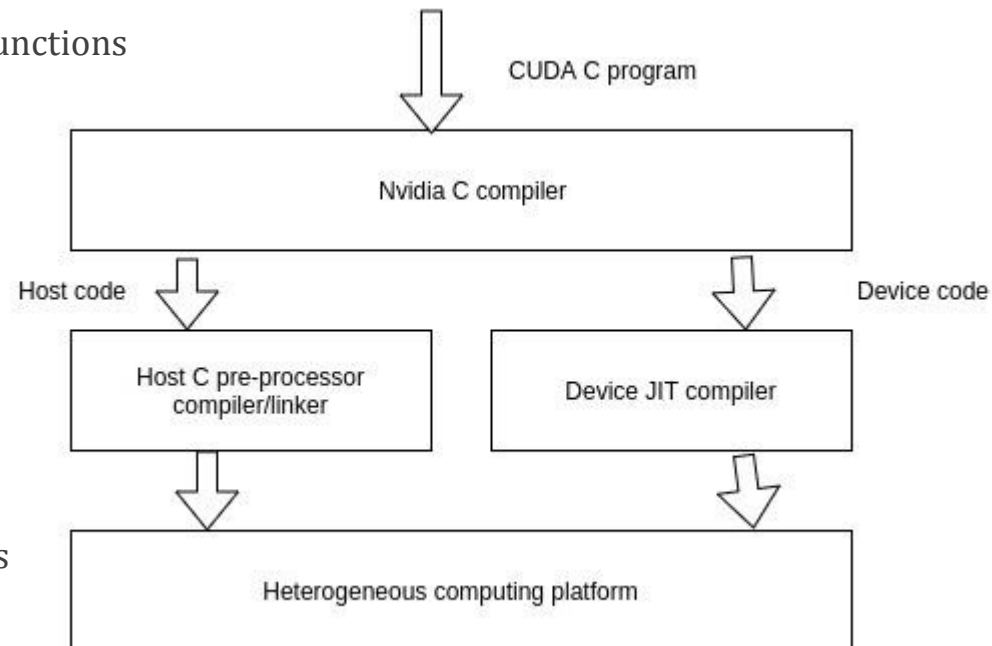
- A compiler to understand API functions

► Host code

- C code - run by CPU
- Compiled by GCC

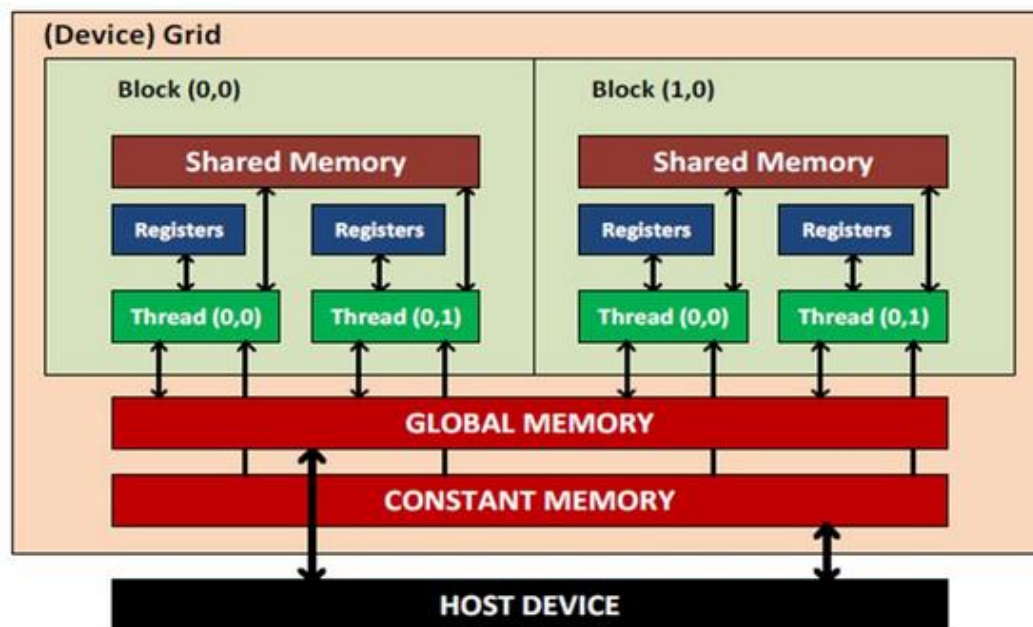
► Device code

- C code - run by GPU
- Special keywords needed
- **Kernels**: data-parallel functions



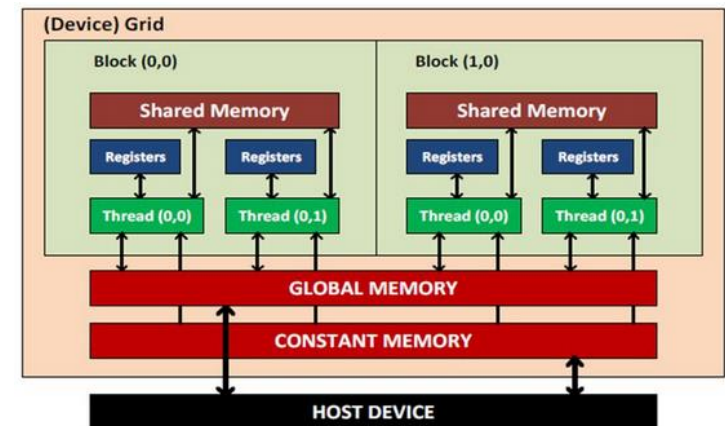
CUDA key concepts

- ▶ Device Global Memory (DRAM)
 - ▶ A typical GPU comes with its own global memory



CUDA key concepts

- ▶ Device Global Memory - API
 - ▶ Allocating memory on the device:
 - ▶ Transfers data from the host to the device memory
 - ▶ Kernel execution on the device:
 - ▶ Transfers the result from the device memory to the host
 - ▶ Free-up allocated memory on the device:
 - ▶ Transfer data to and from the device memory



CUDA key concepts

► CUDA programming keywords:

	EXECUTED ON THE –	CALLABLE FROM –
<code>__device__</code> float function()	GPU (device)	CPU (host)
<code>__global__</code> void function()	CPU (host)	GPU (device)
<code>__host__</code> float function()	GPU (device)	GPU (device)

- `cudaMalloc()` allocate memory on the host
- `cudaFree()` release objects from device memory
- `cudaMemcpy()` memory data transfer

CUDA key concepts

► CUDA programming keywords:

```
void vecAdd(float* A, float* B, float* C,int N) {  
    int size=N*sizeof(float);  
    float *d_A,*d_B,*d_C;  
    cudaMalloc((void**)&d_A,size);
```

```
    cudaMemcpy(d_A,A,size,cudaMemcpyHostToDevice);
```

```
    cudaMemcpy(C,d_C,size,cudaMemcpyDeviceToHost);
```

```
    cudaFree(d_A);
```

```
    cudaFree(d_B);
```

```
    cudaFree(d_C);
```

CUDA key concepts

- ▶ CUDA programming keywords:
 - ▶ blockIdx index of a block in a grid
 - ▶ threadIdx index of a thread in a block
 - ▶ blockDim the dimensions of the grid
 - ▶ blockDim the dimensions of the block
 - ▶ dim3 a data structure (like int) for threads
 - ▶ kernelName<<<#blocks , #threads>>>(parameter1, parameter2, ...)

Hint: threads → blocks → grids.

CUDA installation

- ▶ Requirements (commonly):
 - ▶ A CUDA-enabled NVidia GPU
 - ▶ A supported version of Microsoft Windows
 - ▶ A supported version of Visual Studio
 - ▶ The latest CUDA toolkit

<https://developer.nvidia.com/cuda-downloads>

Select Target Platform ⓘ

Click on the green buttons that describe your target platform. Only supported platforms will be shown.

Operating System	Windows	Linux	Mac OSX		
Architecture ⓘ	x86_64				
Version	10	8.1	7	Server 2016	Server 2012 R2

CUDA installation

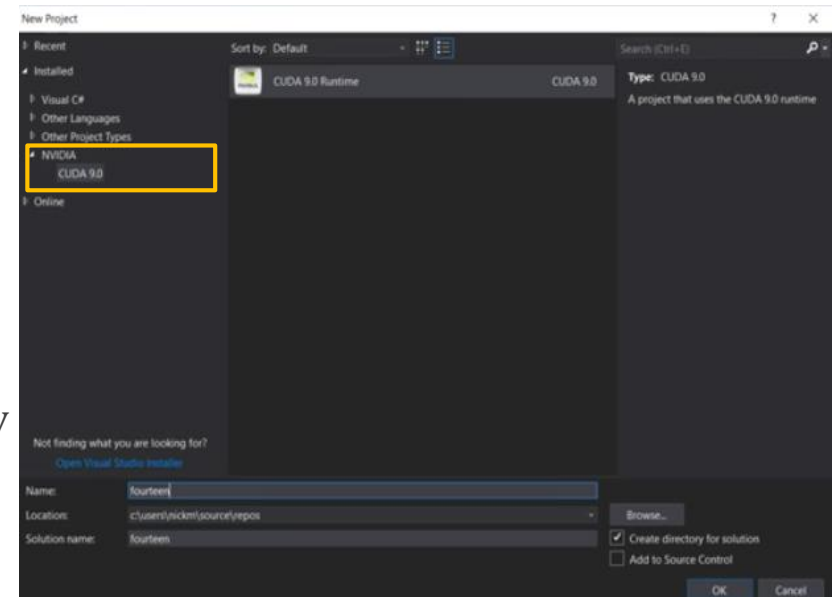
► Requirements (commonly):

Visual Studio Version	Native x86_64 support	X86_32 support on x86_32 (cross)
2017	YES	NO
2015	YES	NO
2015 Community edition	YES	NO
2013	YES	YES
2012	YES	YES
2010	YES	YES

Windows version	Native x86_64 support	X86_32 support on x86_32 (cross)
Windows 10	YES	YES
Windows 8.1	YES	YES
Windows 7	YES	YES
Windows Server 2016	YES	NO
Windows Server 2012 R2	YES	NO

CUDA installation

- ▶ Setting-up Visual Studio for CUDA:
 - ▶ Open Visual Studio → New → Project
 - ▶ Select NVidia from left pane
 - ▶ Creates a project with default codes
 - ▶ Tool set (project properties)
 - ▶ Select suitable platform toolkit
 - ▶ under CUDA C/C++:
 - ▶ Select Common
 - ▶ Set the CUDA Toolkit Custom Directory
 - ▶ Build the project!



CUDA Matrix Multiplication

- ▶ How to store matrices?
 - ▶ Row-major (C, C++, CUDA, ...)
 - ▶ Column-major (Fortran, ...)

M0,0	M0,1	M0,2	M0,3
M1,0	M1,1	M1,2	M1,3
M2,0	M2,1	M2,2	M2,3
M3,0	M3,1	M3,2	M3,3

M0,0	M0,1	M0,2	M0,3	M1,0	M1,1	M1,2	M1,3	M2,0	M2,1	M2,2	M2,3	M3,0	M3,1	M3,2	M3,3
------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------

M0,0	M1,0	M2,0	M3,0	M0,1	M1,1	M2,1	M3,1	M0,2	M1,2	M2,2	M3,2	M0,3	M1,3	M2,3	M3,3
------	------	------	------	------	------	------	------	------	------	------	------	------	------	------	------

CUDA Matrix Multiplication

► Addressing in CUDA

- $\text{element}(x,y) \rightarrow x * \text{width} + y$
- Each element should be mapped to a thread

```
row=blockIdx.x*blockDim.x+threadIdx.x;  
col=blockIdx.y*blockDim.y+threadIdx.y;
```

- All threads in the same block have the same block index (e.g. yellow one)

blockDim.x=4 blockDim.y=1 blockIdx.x=0 threadIdx.x=0/1/2/3

M0,0	M0,1	M0,2	M0,3
M1,0	M1,1	M1,2	M1,3
M2,0	M2,1	M2,2	M2,3
M3,0	M3,1	M3,2	M3,3

CUDA Matrix Multiplication

► Final Kernel:

```
__global__ void simpleMatMulKernell (float* d_M, float* d_N, float* d_P, int width) {  
    int row = blockIdx.y*width+threadIdx.y;  
    int col = blockIdx.x*width+threadIdx.x;  
    if(row<width && col <width) {  
        float product_val = 0;  
        for (int k=0; k<width; k++) {  
            product_val += d_M[row*width+k]*d_N[k*width+col];  
        }  
        d_p[row*width+col] = product_val;  
    }  
}
```

CUDA Matrix Multiplication

```
__global__ void MatrixMulTiled(float *Md, float *Nd, float *Pd, int width) {  
    __shared__ float Mds[width_tile][width_tile];  
    __shared__ float Nds[width_tile][width_tile];  
    int bx = blockIdx.x;  
    int by = blockIdx.y;  
    int tx = threadIdx.x;  
    int ty = threadIdx.y;  
  
    //identify what pd element to work on  
    int row = by * width_tile + ty;  
    int col = bx * width_tile + tx;  
    float product_val = 0;  
    for(int m = 0; m < width/width_tile; m++) {  
  
        //load the tiles  
        Mds[ty][tx] = Md[row*width + (m*width_tile + tx)];  
        Nds[ty][tx] = Nd[col*width + (m*width_tile + ty)];  
        __syncthreads();  
        for(int k=0; k < width_tile; k++) {  
            product_val += Mds[ty][k] * Nds[k][tx];  
        }  
        pd[row][col] = product_val;  
    }  
}
```

CUDA Matrix Multiplication

► Sample:

► Final matrix: d_P (3x3)

► for element (2,1):

► row=2 column=1

2	4	1		4	8	9
8	7	4	×	1	7	0
7	4	9		2	5	4

```
product_val = 0 + d_M[2*3+0] * d_N[0*3+1]
product_val = 0 + d_M[6]*d_N[1] = 0+7*8=56
```

```
product_val = 56 + d_M[2*3+1]*d_N[1*3+1]
product_val = 56 + d_M[7]*d_N[4] = 84
```

```
product_val = 84+d_M[2*3+2]*d_N[2*3+1]
product_val = 84+d_M[8]*d_N[7] = 129
```

Any Questions?



Resources

- ▶ Tutorials-point website (fetched December 2018), see: https://www.tutorialspoint.com/cuda/cuda_performance_considerations.htm
- ▶ D. De Donno, A. Esposito, L. Tarricone and L. Catarinucci, “Introduction to GPU Computing and CUDA Programming: A Case Study on FDTD,” IEEE Antennas and Propagation Magazine, vol. 52, no.3, June 2010.
- ▶ J. Sanders and E. Kandrot, “CUDA by Example,” NVIDIA, 2008.